

# Two-Column Research Note

Introduction. Retrieval quality depends heavily on chunk boundaries. Sentences that are split across chunks lose context, while chunks that are too large dilute relevance scoring. This note evaluates three chunking strategies on a held-out corpus of technical documents.

Results. The sentence-aware strategy outperformed fixed-width chunking by 8 points of recall at  $k=10$ . Overlap of one sentence between adjacent chunks recovered most of the boundary-loss cases without materially increasing token cost.